

The Clay Engine

How we turn 5,000 credits into a pipeline motion that finds the right accounts, writes the right message, and hands reps a ready-to-work list every morning.

The idea in one breath. Clay is an always-on production line for pipeline. It finds the companies worth pursuing, finds the right people inside them, drafts a personalized message for each one, and routes an approved list to a rep. The machine does the manual work. People keep the judgment.

Prepared for **Naveen Thilagan** · Director, Product Strategy

Author **Dallas Andrews** · GTM Engineer · July 2026 · Internal working document

The Clay engine, explained simply

Clay is built from tables that feed other tables, like stations on an assembly line. A record only moves to the next station if it passes a check at the current one. Nothing is manual once it is wired, and the whole line refreshes itself. It runs eight stations deep. *Click a station.*

STATION 1 • SOURCES

Where records come in



STATION 2 • ACCOUNTS

The Golden List



STATION 3 • SIGNALS

Why now



STATION 4 • CONTACTS

The buying committee



STATION 5 • MESSAGE

The draft



STATION 6 • APPROVAL

The human gate



STATION 7 • SEND & TRACK

Out the door, measured



STATION 8 • REP WORKLIST

Who to work today

STATION 1 • SOURCES

Where records come in

Public CMS Star Ratings data for new-logo targets, and our own customer list from Salesforce for the back office. A record only enters if it qualifies, for example a plan sitting under 4.0 stars that is not already a customer.

Every arrow is an automatic Clay action, not a copy-paste. Reference tables (approved claims, the persona rulebook, the message library) sit alongside the line and are read by the stations that need them.

How we protect the 5,000 credits

Clay runs on a fixed budget of credits we spend on enrichment. The plan is built to spend them on the right rows only, show the results, and use that record to renew. Three guardrails keep the spend honest.

GUARDRAIL 1

Never pay twice

A row is only enriched if that field is still empty. Once we have someone's email, we never spend a credit re-finding it.

GUARDRAIL 2

Only enrich what qualifies

Out-of-window and out-of-profile rows are dropped before any enrichment runs. We never pay to research a company we would not pursue.

GUARDRAIL 3

Cheap data first

Low-cost company details first; premium data (intent, hard-to-find emails) only on rows already worth it.

On top of that, a running ledger tracks every batch: credits in, contacts produced, messages sent, qualified replies, meetings. That ledger is the renewal case, told in receipts rather than promises. A checkpoint at 60% of the budget forces a review before we spend the back half.

5,000

Clay credits purchased, team trained, security review done

40+

data sources feeding enrichment per record

60%

spend checkpoint before the back half is released

The first three figures are confirmed on the onboarding site. Every dollar and ROI figure the engine generates is labeled as a modeled estimate until an Intradiem-verified number replaces it.

06 • THE QUICK WIN

The back office inside our own customers

The fastest place to show impact, and it is Mandate 3: 200-plus back-office contacts.

Our front-office customers cannot hand us leads into their own back offices, so we source those contacts ourselves through the same engine, then run campaigns inside accounts where we are already trusted.

WHY IT'S FAST

Warm logos

We are already in the building on the front office. The back office is white space next door, not a cold start.

WHY IT'S FAST

A defined universe

The target functions are known: claims, enrollment, billing, payments, document processing. We point the machine at a list, not the open market.

WHY IT'S FAST

A message that writes itself

"We run your phones, we don't yet see the work behind them" is a clean, honest opener only we can send.

The one gate before we source

The back-office target profile gets ratified with Scott Kemme first. His edits *are* the profile. We review the first fifty contacts together, then the list-build runs against the agreed definition and paces to the Back Office Optimizer launch, not a dump. This keeps the 200 contacts credible instead of a raw count.

07 • THE CHANNEL QUESTION

Where LinkedIn Sales Navigator plugs in

Sales Navigator fits the engine best as a feed into the machine we already have, not as a separate tool a rep works by hand. It plugs in at three points.

PLUG-IN 1

A people source

Its filters (title, seniority, function, geography) feed the Contacts station directly. It is the cleanest way to find the exact buying committee, and it is especially strong for the back-office personas, where public CMS data gives us the account but not the people.

PLUG-IN 2

A signal source

It surfaces job changes, new-in-role moves, and hiring, which are exactly our priority triggers. A new quality leader or a claims-ops hire fires a fresh, timely angle instead of a generic touch.

PLUG-IN 3

A channel

LinkedIn is already the second touch in every sequence. A light automation layer runs the connection, follow, and message steps in step with the email, so the same contact gets one coordinated push.

My recommendation

Bring it in as a **data and signal source first**. That upgrade needs no new behavior from reps and immediately makes the engine's contact and timing quality better. Layer the **outreach automation second**, once deliverability is green and the profile is ratified, so we stay coordinated rather than spraying.

One guardrail to plan for. LinkedIn automation has account-safety limits, so we pace it under daily caps and keep the same human-approval gate we use on email. Same discipline, one more channel.

This is the path I'd take. I'll scope the specific tool and its cost so we're deciding on a real number, not a maybe.

08 • THE OPEN THREADS

Where I'd value your read

The engine and the messaging are built and ready to stand up in our Clay workspace. A handful of things sit with you rather than with the build, and I'd rather shape them with you than assume. Here is where your read moves this forward.

1 • What "engine-sourced" means, defined together.

A shared definition (and the three Salesforce tracking fields behind it) keeps the ledger and the credit story a common scoreboard rather than something to debate later.

2 • The back-office profile, shaped with Scott Kemme.

His edits become the sourcing spec, and I'd build the 200 contacts from there. Whenever that conversation makes sense on your side, I'm ready to run with the outcome.

3 • The deliverability bar before we scale.

I'll prove on a small slice and watch inbox health first. Setting the threshold together is what keeps a fast start from costing us the domain.

Where this is heading. Every one of the new DWO products, Back Office Optimizer and Queue Optimizer first, needs pipeline aimed at the right buyer at the right moment. This engine is that front end. It's built and ready to stand up in our Clay workspace, aimed at the Star Ratings play first, with the back-office motion the next thing we point it at. It's designed to get sharper with every campaign rather than needing a rebuild. That compounding is the point, and it is the part that keeps paying off as the product line grows.
