

— VALIDATION FINDINGS • GTM ENGINEERING

ACD Detection: Validation Findings

Internal briefing prepared for **Naveen Thilagan** • GTM Engineering • July 2026

A quick test of a big idea, before we build a play on top of it. This is what the data says, where it holds, and where it doesn't, so we can decide together how far to lean on it.

01 THE IDEA BEING TESTED

The idea came out of a strategy call with Josh, our lead solutions consultant. Intradiem integrates with a customer's ACD and WFM, so the sharpest outreach personalization we could run is to detect a prospect's ACD and lead with a proof story from a customer on that same stack. Same platform, same integration story, a peer who already made it work. The question this test answers is narrow and load-bearing: **can we detect a company's ACD accurately enough to personalize at scale?**

02 HOW WE TESTED IT

Built a Clay detection workflow (domain-find, then two tech-detection providers, HG Insights and BuyerCaddy, run blind, then a match-score step) and ran it against **39 existing Intradiem customers whose true ACD we already know** from the internal customer export. Known answers on every row, so the test grades itself. Built and run in one day.

17 / 39

HG INSIGHTS
CORRECT • 44% RAW

~50%

AFTER DEDUPING
SHARED-DOMAIN SUB-
ENTITIES

0 / 39

BUYERCADDY CORRECT
• NOT YET
CONCLUSIVE

~486 cr

FULL VALIDATION
COST • INSIDE
MONTHLY BUDGET

1 day

BUILD AND RUN, END
TO END

03 THE HEADLINE RESULT

HG INSIGHTS

17 / 39


44% raw, roughly 50% after deduping. Several Kaiser sub-entities share one domain, so a single miss counts against several rows; correcting for that lifts the rate to about half. This is the provider carrying the signal.

BUYERCADDY

0 / 39

Returned "not in use" for every candidate on every account. Likely blind to back-office ACD, but that's a hypothesis. **Pending a catalog check before we call it, so this is not yet conclusive.**

04 ACCURACY BY ACD VENDOR · HG Insights



Read it by vendor, not as one blended number. Cisco and Five9 detect cleanly; the modern-cloud and legacy names (Amazon Connect, Avaya) run soft; NICE has too small a sample to read at all. The play's confidence isn't uniform, it's vendor-specific.

05 THE NUANCE, KEPT IN VIEW

THE 44% IS A FLOOR

Four accounts failed domain resolution because of a fixable naming artifact in the test set, not because detection failed. Correct that and the true rate likely lands in the **low-50s**. The headline understates the ceiling, so read it as the bottom of the range, not the middle.

BIG ENTERPRISES RUN SEVERAL

On accounts like **UnitedHealth, Aetna, Home Depot, and Kaiser**, detection returns multiple real ACDs, because they genuinely run several across business units. The tool isn't wrong there; "which ACD" simply isn't a single answer for a large enterprise, and a match-story play has to account for that.

06 THE VERDICT

WHERE IT STANDS

ACD detection isn't reliable enough to be the backbone of a blanket, automated "match-story-to-ACD" play across cold prospects. A confidently wrong guess about a prospect's own stack would damage the exact credibility the play is built to earn. **It is genuinely usable two ways, though, and both are worth having.**

01 • THE NARROW HIGH-CONFIDENCE PLAY

Fire the ACD-matched message only on **Cisco and Five9 detections**, where accuracy is 80 to 100%. Small, sharp, and safe to automate.

02 • AN INTERNAL RESEARCH SIGNAL

Use detection as a **prioritization and research input a human verifies before send**, not an auto-personalization trigger. Speeds the work without staking credibility on it.

PROVISIONAL · FOR US TO DECIDE

Near term, I'd lean toward running the high-confidence subset now (Cisco and Five9) while treating everything else as a human-verified research signal. That captures the value that's real today without betting the play's credibility on a guess.

The bigger question is the data source. **We already hold the true ACD for every customer;** the gap is net-new logos, where detection is exactly this shaky. The real unlock is a stronger ACD data source for prospects, and whether that's worth an investment before we scale the play feels like a call worth making together rather than one I'd make on my own read.

~486 credits

The full validation ran on about 486 Clay credits, well inside the monthly operating budget. **We answered a flagship-play question with receipts in a day, and found its limit cheaply, before building on top of it.**