

- INTERNAL ONLY. DESCRIBES THE OUTBOUND ENGINE AND NAMES PROSPECT EXAMPLES. NOT FOR EXTERNAL OR PROSPECT-FACING USE.

— MESSAGING ENGINE UPDATE • GTM ENGINEERING

The Stars Outreach Engine: What Changed, and What Earns a Reply

Prepared for Naveen Thilagan and Nathan Belfield • GTM Engineering • July 22 2026

Nate's notes and everything from the 1:1 are folded into the copy engine. This is what the messaging does now, why every email clears the bar before it sends, and an honest read on what actually earns the replies, so we judge it the right way.

01 WHAT YOU FLAGGED, AND WHAT THE ENGINE DOES NOW

BEFORE

- Opened by restating the reader's own star number
- Named specific **contract IDs** we couldn't be sure they owned
- Canned lines like "**let me make this concrete**" on every lead
- Read dense and a little AI
- Said little about what Intradiem actually does

NOW

- Opens with an **insight**; their number lands as support
- **No contract IDs**. Everything is payer-level
- Canned openers banned; each email is built from that contact
- Warm, plain, short paragraphs, in Nate's register
- One plain product line, framed as our **Dynamic Workforce Orchestration** platform, plus one verified proof

5 touches

JUDGE THE
SEQUENCE, NOT ONE
EMAIL

3 + human

AUTOMATED
CRITICS, THEN A
PERSON, BEFORE
ANY SEND

1 proof

HUMANA, VERIFIED
& CLEARED FOR
PROSPECT USE

0 IDs

CONTRACT IDS IN
COPY, PAYER-LEVEL
ONLY

46 leads

REGENERATING ON
THE NEW VERSION

02 HOW THE COPY READS NOW

Every email opens with a point of view, not the prospect's own score, because a Stars or Finance leader knows their gap cold and being told it reads as homework not done. The story is one clear chain, stated as upside: **Intradiem improves the customer-service and operational metrics, and the service Star measures rise as a result.** We never claim to move a measure directly, and we never touch clinical. One plain line names what our Dynamic Workforce Orchestration platform does (shift and reprioritize work across queues in real time to hold service levels), and one verified customer result carries it (Humana's handle-time improvement), framed operationally, never as a Stars claim. The close gives the reader a specific reason worth 15 minutes, not a generic "let's connect."

The five touches are written as one building story to someone who hasn't replied, not four sales pitches: a first note that earns attention, a softer second angle, the stakes, the one-page breakdown, then a gracious close. Each is the next natural note from a sharp peer.

03 WHY YOU CAN TRUST WHAT GOES OUT

Nothing sends on the model's say-so. Each email clears a stack of automatic checks and then a human:

- 01 **Figures check.** Every number must trace to public CMS data or our one cleared proof point. Anything else fails.
- 02 **Voice check.** Fails cold or robotic copy, canned openers, and any line that reads as AI-written.
- 03 **Full five-email pre-send audit.** Reads the whole sequence together before any of it can send.
- 04 **Claims gate.** Only Humana and public CMS math are quotable. No other customer, no clinical claim, no contract ID.
- 05 **Human approval.** Nathan reads and sends from his own mailbox. Campaigns are paused now on purpose while we regenerate all 46 on the new version.

WHAT I'D ACTUALLY TELL YOU

Would I send these? Yes. They're well above typical cold outbound: they open with a point of view instead of the prospect's own number, they're specific, they concede the clinical side (which buys credibility), the proof is real and carefully framed, and the CTA gives a reason worth 15 minutes. I stand behind the craft.

Would I promise replies? No, and I'd be suspicious of anyone who did. Cold reply rates sit in the low single digits even for excellent copy. What actually earns the replies is three things working together, and only one of them is the writing.

NUMBER
ACCURACY

The real make-or-break for these readers. A VP of Stars knows her weighted average cold. If a figure is off, the email doesn't just miss, it burns trust and files us as "didn't do the homework." **Every number is verified against the CMS data in the table before send.** That single discipline moves reply odds more than any wording tweak.

DELIVERABILITY

Replies require landing in the primary inbox from a warmed mailbox. Great copy in the spam folder gets zero. **That gate is separate from the writing** and we treat it as its own thing.

SEQUENCE +
PHONE

It takes eight to ten touches for most people to reply, and rarely a single email. These first notes are strong openers, but the full five-email arc and Nate's calls off the sourced numbers are what convert. **Judge the sequence, not the single send.**

05 WHAT'S NEXT, THREADS TO SHAPE TOGETHER

PROVISIONAL • FOR US TO DECIDE

Nate's own best-performing emails get loaded as the engine's voice anchor, so it aligns to what already works rather than guessing. Then we grade a fresh batch together, keep the ones that land, and the next seven-company wave runs on this same version so the comparison is clean.

One call worth making together: confirming the enriched data points we lead with are the ones worth trusting, so the personalization stays credible as we scale. Maya is the right person to pressure-test the data side with.

Nothing sends until the numbers are verified against the table and you and Nate sign off on a batch. Then we regenerate all 46 and turn the sequence back on.

Dallas Andrews · GTM Engineer, with Naveen Thilagan & Nathan Belfield