



GTM ENGINE BRIEFING · JOHN · TOM · CHRIS · JENN · JULY 2026

The Stars motion has a full-time engine behind it, built to run weekly.

Where the build stands this week, the workbook itself, and where the engine goes after Star Ratings. One decision at the end: the green light to start.

Dallas Andrews · GTM Engineer · Product Strategy, with Naveen Thilagan

The Stars window moves weekly. The motion has to match it.

Signals fire continuously

CMS releases, filings and earnings language, quality leadership moves, hiring clusters



Every seat needs researched, personal outreach

Roughly five committee seats per parent, each with its own angle, never a template blast



At a pace no calendar sets alone

The mechanical work runs at machine speed; judgment and conversations stay human

32

TARGET PARENTS, MOST
BONUS REVENUE AT
RISK

143

COMMITTEE CONTACTS
TO WORK

5

TOUCHES PER SEAT,
MULTI-CHANNEL

~700

RESEARCHED, PERSONAL
TOUCHES PER FULL
PASS

The accounts in that window carry nine-figure Stars money.

Every parent in the target universe is scored against the 2026 CMS release, with forgone quality-bonus dollars estimated from public CMS data and enrollment. Estimates, clearly labeled; each plan knows its exact picture better than we do. Two examples already graded and staffed with committee contacts:

\$165.8M est.

Centene estimated forgone QBP, 21 cliff-edge contracts on the 2026 cycle. A confirmed non-customer, in scope for new logo.

\$61.2M est.

Medica estimated forgone QBP, 3 cliff-edge contracts on the 2026 cycle. Committee contacts already in the table.

The window in plain terms: what these plans do in the 2026 and 2027 measurement years is what the 2028 Star Ratings will grade, and the call-center measures we credibly move still count in that window before CMS retires them. Forgone bonus is forgone revenue, dollars that flow straight into the plan's top line when a contract clears 4.0.

Source: CMS Oct 2025 release + enrollment, KFF/MedPAC methodology. Estimate label travels with every number.

A live engine in Clay, already scored, staffed, and gated.

CONTRACTS

307

2026 CMS Star Ratings contracts scored and graded live on native formulas

CONTRACTS

98 / 32

Tier A+B contracts across 32 parent organizations

PEOPLE

143 / 32

Verified quality and finance decision-makers across all 32 target parents, sourced, deduped, and validated; rows at customer parents held for the expansion lane

PEOPLE

3

Personas pulled: Stars/Quality leaders, Medicare-segment finance including Chief Actuaries, and contact-center and member-services operations


LOCKED OUT OF EVERY SEND

Current customers are locked out of every send. The lock lives in the engine itself: the sync that feeds each campaign refuses any row at a customer parent, reconciled against the Salesforce active-customer report. Expansion runs its own lane through the account managers. The two motions cannot collide.

GATED BY DEFAULT

Wave 1 is staged: two persona campaigns holding 38 leads in draft, 23 Stars/Quality and 15 Finance, with a 15-lead longer-horizon lane staged behind them. Every lead is through the engine's approval gate and a per-lead human review before launch. Nothing has gone out, and nothing can, until a person turns the key. Zero sends is the gates working, not the engine waiting.

From 307 contracts to 38 staged leads. Every cut on the record.



No Intradiem metric on any row. Every figure walks back to the CMS release or the Salesforce customer report, and every cut above is a filter you can audit.

Total build to date: about 820 credits. Roughly \$57 of research.

DATE	ACTION	CREDITS	NOTE
07-08	Full 98-row Tier A+B universe imported, formulas verified live	0	Imports and formulas are free by design
07-08	Accounts wired to Contacts; grades recompute live	0	Internal lookups, meter-verified at zero
07-08	Two-persona committee pull + work-email waterfall, 94 contacts	142.3	Estimated 94, actual 142.3, meter-verified
07-09	Email validation on the 78 emailed rows	7.8	Pre-estimated 6–10; 0.1/row; 67 deliverable
07-09	Ops-persona gap-fill: 43 contacts, all 27 parents now covered	~48	Import free; email waterfall pre-quoted ~48.4
07-09	Validation on the new rows	5.9	0.1/row; ~106 deliverable total
07-10	Salesforce active-customer file reconciled, 101 accounts; customer lock wired into the engine	0	The sync refuses every row at a customer parent
07-10	Wave 1 approvals: HOLD	0	Nothing sent; no approval box ever checked
07-11	Wave 1 campaigns built and loaded, sequences from the approved set	0	Native campaign build; holds in draft behind review
07-12	First live signal sweep + adversarial audit of the whole build	~139	Eligible-only, meter-verified; findings fixed same day
07-13	Both load lists rebuilt to full addressable width, 34 + 24 leads	0	Every cut on the record; campaigns hold in draft
07-13	Coverage gap-fill: 6 new committee seats sourced and validated	~35	Meter-verified; the table now holds 143 rows
07-14	Per-lead personalization layer + full audit-and-regen hardening waves, all 38 send-surface drafts	~324	Meter-verified 378 → 702; every draft re-audited before any sync
07-15	Wave 1 staged at full width, 23 + 15 leads in draft; longer-horizon lane built and staged, 15 more	~115	Meter-verified; one lane per contact, three draft campaigns

~820 cr

ENTIRE BUILD TO DATE
· ABOUT \$57

+1,150–1,500 cr

REMAINING COMMITTEE
EXPANSION, PRE-
ESTIMATE

≤2,300 cr

FULL STARS MOTION ·
ABOUT \$160

Traction-led

SPEND FOLLOWS WHAT
REPLIES · MOTIONS
AUDITED CONTINUOUSLY

Credits are Clay's research unit, not dollars: about 7 cents each. Append-only ledger, one line per action; nothing runs without a logged pre-estimate and an explicit go.

Planning ranges, replaced by our own numbers after Wave 1.

WAVE 1 • 38 CONTACTS • THE CALIBRATION WAVE

No result targets on this wave, on purpose. 38 contacts is deliberately too small to judge the motion; a single reply moves the rate by 3 points. Its job is real deliverability, reply, and persona data, so the ranges below become our numbers instead of industry assumptions before we put volume behind them.

TRACTION SETS THE BUDGET

The ranges are planning assumptions, set conservative. As replies show which personas and messages earn attention, **the motion doubles down where the market answers:** every lane is continuously audited, spend follows the evidence, and if the results say push harder, pushing harder is the easy decision.

143

CONTACTS WORKED

6–11

REPLIES • 4–8% RATE

3–7

QUALIFIED
CONVERSATIONS

2–4

MEETINGS TO SALES

1–2

OPPORTUNITIES
SOURCED

Attribution set on day one. Reported where you already look.

REPLY → QUALIFIED → MEETING

Every contact is tracked from first send through reply, qualified conversation, and meeting booked. Qualified means a target persona engaging on the problem; auto-replies, referrals, and polite declines do not count. Attribution is set the day a contact enters the system, not reconstructed later.

THE ACTIVITY LAYER

Sends, deliverability, and reply rates by persona ride the same readout, so a quiet week reads as a targeting problem or a deliverability problem, never a mystery.

COST PER QUALIFIED REPLY

Every enrichment credit is logged against the motion that spent it, so spend ties directly to output. The unit economics of the engine, reviewable any week.

SOURCED VS. EXISTING, NEVER BLENDED

Engine-sourced replies, meetings, and pipeline report separately from anything already in motion, and estimated opportunity never blends with realized results. The engine only gets credit for what it actually created.

Reporting lives where you already plan: engine results feed Salesforce directly, with actuals against these ranges in a dashboard you can open any time, shaped to how pipeline council and the H1 and H2 reviews already look at pipeline. **Nothing new to log into.**

The list keeps itself alive. Half of that costs nothing.

FREE · COMPUTED FROM DATA ALREADY IN THE ENGINE

The cliff-edge signal. Fires points when a contract sits just under 4.0 and self-cleans to zero the moment it clears the line, so an account that solves its problem drops off the list the same day. Intent rollup, fire dates, and the tier logic (Tier 1 fires outreach, Tier 2 raises priority) run the same way: native formulas, zero credits, the most important signal for the cliff play.

PAID · ENRICHMENTS, PRE-ESTIMATED PER RUN

The outside world. Stars language in SEC filings, earnings-call mentions, new quality leadership, quality hiring clusters, backlog news. Each is an enrichment per parent, roughly 2.5 credits per signal, so a full five-signal sweep across all 32 parents runs about 400 credits, logged in the ledger before it fires like everything else.

Monthly

FULL FIVE-SIGNAL
SWEEP, ~400 CREDITS

Weekly

TOP-TIER DELTA ON
FAST MOVERS, ~100
CREDITS

Event-driven

FULL SWEEP ON
EARNINGS WEEKS + THE
OCTOBER CMS RELEASE

~800 / mo

STEADY STATE TODAY ·
SCALES WITH TRACTION

Per-provider costs confirmed against Clay's current pricing before anything fires. The cadence is a dial, not a commitment; moving it is a config change.

Three things to see in the system.

IN THE WORKBOOK

THE LIVING LIST

Scoring reacts to data: contacts land and the account's grade recomputes, and the cliff-edge signal zeroes itself the moment a contract clears 4.0, so a plan that solves its problem drops off the list the same day. Not a report, a living list.

IN THE WORKBOOK

THE GATE RAIL

Every row carries its whole gate on one screen: the automated critic's verdict, the human approval, the send-ready state. The sync that feeds a campaign refuses any row missing one of them, and all three campaigns hold in draft until a person deliberately launches.

IN THE WORKBOOK

COLLISION CONTROL

One checkbox: a rep claims a contact, the engine stands down on that person and their whole company, outranking any approval. A reply stops that contact's sequence on its own; the company-wide stand-down is one click behind it.

The tables stay open the whole session, and anything you want to see is one click away. And whatever this room decides it wants eyes on becomes a standing readout that finds you; the engine's job is to bring you answers, not another tool to log into.

A wave a week without new headcount. Every number traces to its source.

A WEEKLY MOTION, BY DEFAULT

Templates, scoring, gates, and the ledger operate as one repeatable cycle: a full wave ships every week, and human judgment concentrates in the approval gate, where it decides what reaches a prospect instead of being spent on mechanical work. The wave itself spends nothing; sending is free, credits price research, and that research is already inside the steady state.

PIPELINE YOU CAN TRACE TO ITS SOURCE

Every contact is stamped with its source motion and sourced date at creation. Surfaced opportunity reports separately from realized, never blended, so when a number reaches your forecast you can walk it back to the row that produced it.

Marketing fires on the same signal, at the same moment.

TRIGGERED CAMPAIGNS

One account signal produces three synchronized outputs: a campaign brief for marketing, an alert for the owning seller, and a gated outreach draft. Marketing and sales act on the same trigger at the same time instead of finding out from each other later.

THE CONTENT ENGINE

One source asset, a webinar, a customer story, a conference talk, becomes a coordinated drop: blog post, LinkedIn set, nurture sequence, sales one-pager, all consistent and on brand in one pass. Volume without the team growing.

This engine is built as a technology partner to marketing, not a parallel motion, and the near-term prize is the story vault: **two decades of customers who stayed, with story after story still to tell.** Any one of them, run through the engine, becomes a full coordinated drop with every claim tied to a verified source, and the engine hands marketing more surface area on triggers no one could watch manually.

Meetings arrive pre-scored. Reps never collide with the engine.

COMMITTEE CONTEXT ATTACHED

Before the first call, the seller has the account grade, tier, the estimated Stars dollars at stake, and which committee members are already engaged. The meeting starts at the second conversation, not the first.

ONE CONTACT, ONE OWNER

A reply stops that contact's sequence on the spot, and one claim checkbox stands the engine down on that person and their whole company. From that moment the contact belongs to a human. No double-touches, ever, by construction.

October re-arms the engine. Stars is only its first motion.

CMS re-rates every contract each October. I designed the engine for that moment: the release lands, the universe re-pulls, every grade recomputes, and the motion re-arms behind the same gates, with a fresh and urgent reason to be in every one of these rooms. That is why this is an engine and not a campaign; the flywheel turns every year. And the same machinery is already pointed at what comes next:

BUILT · LIVE IN CLAY

BACK OFFICE

A 101-account target universe with 59 sourced contacts, built behind Stars. Ready as the quick follow-up whenever this room wants it.

ON THE SHELF

THE SIGNAL LAYER

The daily scan across all 32 parents, priced earlier in this deck. Filings, leadership moves, and hiring clusters become triggers the day we turn the dial.

ON THE SHELF

INSTALL-BASE EXPANSION

The rows at customer parents are already sourced and held. The same engine runs the expansion lane through the account managers when we choose to arm it.

ON THE SHELF

COMPETITIVE DISPLACEMENT

The same universe, scored on a competitor trigger instead of a Stars trigger. A new motion is a config change, not a rebuild.

Every sector splits into motions of its own.

Each pill is a runnable motion: its own target universe, signal set, persona lanes, and entry story, all behind the same gates. Star Ratings is one pill on this map, live today, proving the pattern for every other one.

SECTOR	BACK OFFICE OPTIMIZER We replace spreadsheets, not an incumbent	QUEUE OPTIMIZER Queue membership decisions, run on a cadence
Insurance UK · US	Home & property claims Auto & motor claims Life & pensions Health / PMI Pet Travel Commercial & specialty	Home FNOL lines Auto FNOL lines Cat event surge Renewals seasonality
Finance UK · US	Mortgage servicing Lending ops Cards & disputes Retail banking ops Wealth ops	Consent order banks Collections Card services Consumer Duty complaints
Healthcare US	Payer claims ops Revenue cycle Appeals & grievances desks	● STAR RATINGS · LIVE Member services lines Provider lines Sub-4.0 QBP parents, call-center measures, 2028 window. 32 parents · 143 contacts · Wave 1 staged.
Retail UK · US	Order management Returns Fraud review	Peak season surge Order & delivery support
Utilities UK · US	Energy billing Water billing Complaints handling	Outage surge Storm & event response Billing lines

37

MOTIONS ON THIS MAP,
TWO PRODUCTS

1 live

STAR RATINGS, WAVE 1
STAGED

Same gates

SCORED, STAFFED,
HUMAN APPROVED

UK · US

SAME ENGINE, BOTH
MARKETS

**We run GTM the way our
product runs operations:
real-time signals,
automated action, a human
gate before anything
irreversible.**

I sit between the systems so one signal becomes synchronized action across marketing and sales at the same moment. The engine is the proof point of our own thesis, running inside our own walls.

The engine has everything it needs. One decision: the start.

01

DATA · ALREADY HANDLED

Salesforce data already flows into the company data hub AI Enablement runs, so the engine reads it there, exclusion lists and account owners included, without touching Salesforce's API limits. Ad-hoc reporting runs through the Greenlight Salesforce agents. No new integrations needed.

02

TOOLING · ALREADY HANDLED

The judgment layer runs on Claude and the seller-facing skills are staged for Greenlight, all inside the stack AI Enablement already runs. Nothing in this meeting waits on infrastructure.

03

THE GREEN LIGHT

Approve the start. Wave 1 is staged: 38 leads across two persona lanes, a longer-horizon lane behind them, every one through the approval gate and holding in draft. On the yes, the wave ships behind the deliverability gate and results feed a Salesforce dashboard you can open any time. Your read on which accounts go first is what shapes it.

This is a head start, not a finished product. The build compounds weekly, and the continuous engineering is the role. Everything upstream is done and gated; the first wave waits only on the judgment in this room.